

# Market Basket Analysis

## R Markdown

Loading the required libraries

```
library(arules)
library(arulesViz)
```

Let us begin. Loading the data set. Here we are not loading using read.csv() function. Why? If you take a look into the data set out data set does not have variables in it. So this might create problem when we load. The read.transactions() function changes the dataset into a sparse matrix. It makes each row represent a transaction and creates columns for each item that a customer might purchase. Electronidex sells 125 items, so the sparse matrix creates 125 columns. It also changes the data to binary. (1=item purchased in that transaction OR 0=no purchase.)

```
TransactionDataSet <-suppressMessages( read.transactions("ElectronidexTransactions2017.csv",format = c(
summary(TransactionDataSet)
```

```
## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 125 columns (items) and a density of 0.03506172
##
## most frequent items:
##               iMac               HP Laptop CYBERPOWER Gamer Desktop
##               2519               1909               1809
##           Apple Earpods       Apple MacBook Air               (Other)
##               1715               1530               33622
##
## element (itemset/transaction) length distribution:
## sizes
##    0    1    2    3    4    5    6    7    8    9   10   11   12   13   14   15
##    2 2163 1647 1294 1021  856  646  540  439  353  247  171  119   77   72   56
##   16   17   18   19   20   21   22   23   25   26   27   29   30
##   41   26   20   10   10   10    5    3    1    1    3    1    1
##
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.000  2.000   3.000   4.383   6.000  30.000
##
## includes extended item information - examples:
##                                labels
## 1 1TB Portable External Hard Drive
## 2 2TB Portable External Hard Drive
## 3                   3-Button Mouse
```

To view data set as a whole use “inspect (DatasetName)”. This is gonna take a lot of time.

```
inspect (TransactionDataSet[1:5])
```

```
##      items
```

```
## [1] {Acer Aspire,
##      Belkin Mouse Pad,
##      Brother Printer Toner,
##      VGA Monitor Cable}
## [2] {Apple Wireless Keyboard,
##      Dell Desktop,
##      Lenovo Desktop Computer}
## [3] {iMac}
## [4] {Acer Desktop,
##      Intel Desktop,
##      Lenovo Desktop Computer,
##      XIBERIA Gaming Headset}
## [5] {ASUS Desktop,
##      Epson Black Ink,
##      HP Laptop,
##      iMac}
```

```
LIST(TransactionDataSet[1:5])#Lists the transactions by conversion (LIST must be capitalized)
```

```
## [[1]]
## [1] "Acer Aspire"          "Belkin Mouse Pad"      "Brother Printer Toner"
## [4] "VGA Monitor Cable"
##
## [[2]]
## [1] "Apple Wireless Keyboard" "Dell Desktop"
## [3] "Lenovo Desktop Computer"
##
## [[3]]
## [1] "iMac"
##
## [[4]]
## [1] "Acer Desktop"          "Intel Desktop"
## [3] "Lenovo Desktop Computer" "XIBERIA Gaming Headset"
##
## [[5]]
## [1] "ASUS Desktop"      "Epson Black Ink" "HP Laptop"      "iMac"
```

```
length(TransactionDataSet) # length of transaction
```

```
## [1] 9835
```

```
size(TransactionDataSet[1:10]) #No:of items per transaction upto the 10th row
```


```
## [1] 4 3 1 4 4 5 1 5 1 2
```

```
length(itemLabels(TransactionDataSet))# To see the length of item labels.. or lets say the list od labe
```

```
## [1] 125
```

There is total of 125 items & our data set has 9835 transactions

```
itemFrequencyPlot(TransactionDataSet, topN =10, col = rainbow(4),type = "relative")
```




Let us find the list of items which are less popular. The below shows the list of 6 items with less sales volume:

|    |                                    |                                  |
|----|------------------------------------|----------------------------------|
| ## | Logitech Wireless Keyboard         | VGA Monitor Cable                |
| ## | 0.002236909                        | 0.002236909                      |
| ## | Panasonic On-Ear Stereo Headphones | 1TB Portable External Hard Drive |
| ## | 0.002338587                        | 0.002745297                      |
| ## | Canon Ink                          | Logitech Stereo Headset          |
| ## | 0.002745297                        | 0.003050330                      |


In addition to looking at the items, it's also possible to visualize the entire sparse matrix. To do so, use the `image()` function. The command to display the sparse matrix for the first 10 transactions is as follows:

```
image(TransactionDataSet[1:10],
      xlab = "Items (Columns)",
      ylab = "Transactions (Rows)")
```



this visualization will not be as useful for extremely large transaction databases, because the cells will be too small to discern. Still, by combining it with the `sample()` function, you can view the sparse matrix for a randomly sampled set of transactions. The command to create random selection of 100 transactions is as follows:

```
image(sample(TransactionDataSet,100))
```



From the graph we can understand that there is some popular items in the store as few columns seem fairly heavily populated. But overall, the distribution of dots seems fairly random.

## Apriorifunction

```
rules1 <- apriori(TransactionDataSet,parameter = list(supp = 0.005, conf =0.6,minlen = 1,maxlen=10,target

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.6   0.1   1 none FALSE                TRUE     5   0.005    1
## maxlen target ext
##          10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 49
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[125 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [109 item(s)] done [0.00s].
## creating transaction tree ... done [0.01s].
## checking subsets of size 1 2 3 4 5 done [0.01s].
```

```
## writing ... [28 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules1
```

```
## set of 28 rules
```

```
rules1 <- sort(rules1, by = 'lift')
```

These parameters are requesting that the rules cover 10% of the transactions and are 80% correct.

To View the rule

```
inspect(rules1[1:5])
```

|        | lhs  | rhs            | support     | confidence | coverage    | lift     | count |
|--------|--|----------------|-------------|------------|-------------|----------|-------|
| ## [1] | {Acer Aspire,<br>Dell Desktop,<br>ViewSonic Monitor}             | => {HP Laptop} | 0.005287239 | 0.8125000  | 0.006507372 | 4.185928 | 52    |
| ## [2] | {Acer Aspire,<br>iMac,<br>ViewSonic Monitor}                     | => {HP Laptop} | 0.006202339 | 0.6630435  | 0.009354347 | 3.415942 | 61    |
| ## [3] | {Acer Desktop,<br>iMac,<br>ViewSonic Monitor}                    | => {HP Laptop} | 0.006405694 | 0.6363636  | 0.010066090 | 3.278489 | 63    |
| ## [4] | {Dell Desktop,<br>Lenovo Desktop Computer,<br>ViewSonic Monitor} | => {HP Laptop} | 0.006202339 | 0.6224490  | 0.009964413 | 3.206802 | 61    |
| ## [5] | {Computer Game,<br>ViewSonic Monitor}                            | => {HP Laptop} | 0.007422471 | 0.6186441  | 0.011997966 | 3.187200 | 73    |

```
#str(rules_df)
```

Receiving 0 rules means that you will need to experiment with the Support and Confidence values.

Now we recieved : set of 28 rules

When you're experimenting keep in mind:

1. If these values are too high, you will receive no rules or non-helpful rules.
2. If these values are too low, your computational time/memory will suffer, or you'll receive too many rules.
3. To get 'strong' rules, increase the value of 'conf' parameter.

Evaluating & taking a deep look

```
summary(rules1)
```

```
## set of 28 rules
```

```
##
```

```
## rule length distribution (lhs + rhs):sizes
```

```
## 3 4
```

```
## 17 11
```

```
##
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max.
```

```
## 3.000 3.000 3.000 3.393 4.000 4.000
```

```
##
```

```
## summary of quality measures:
```

|         | support   | confidence   | coverage       | lift        |
|---------|-----------|--------------|----------------|-------------|
| ## Min. | :0.005084 | Min. :0.6000 | Min. :0.006507 | Min. :2.343 |

```
## 1st Qu.:0.005465 1st Qu.:0.6124 1st Qu.:0.008948 1st Qu.:2.423
## Median :0.006355 Median :0.6321 Median :0.009964 Median :2.536
## Mean :0.006758 Mean :0.6460 Mean :0.010582 Mean :2.725
## 3rd Qu.:0.007550 3rd Qu.:0.6648 3rd Qu.:0.012125 3rd Qu.:2.940
## Max. :0.010778 Max. :0.8125 Max. :0.017895 Max. :4.186
## count
## Min. : 50.00
## 1st Qu.: 53.75
## Median : 62.50
## Mean : 66.46
## 3rd Qu.: 74.25
## Max. :106.00
##
## mining info:
## data ntransactions support confidence
## TransactionDataSet 9835 0.005 0.6
```

The summary of the rules gives us some very interesting information: 1. The number of rules: 28. 2. The distribution of rules by length: a length of 3 items has the most rules. 3. The summary of quality measures: ranges of support, confidence, and lift. 4. The information on data mining: total data mined, and the minimum parameters we set earlier.

Removing the redundant


```
table(is.redundant(rules1))
```

```
##
## FALSE
## 28
```

Plotting the 10 rules

```
topRules <- rules1[1:5]
plot(rules1)
```

```
## To reduce overplotting, jitter is added! Use jitter = 0 to prevent jitter.
```




```
plot(rules1[1:3], method = "graph", control = list(type= "items"))  
## Warning: Unknown control parameters: type
```




### Graph for 3 rules

size: support (0.005 – 0.006)  
color: lift (3.278 – 4.186)



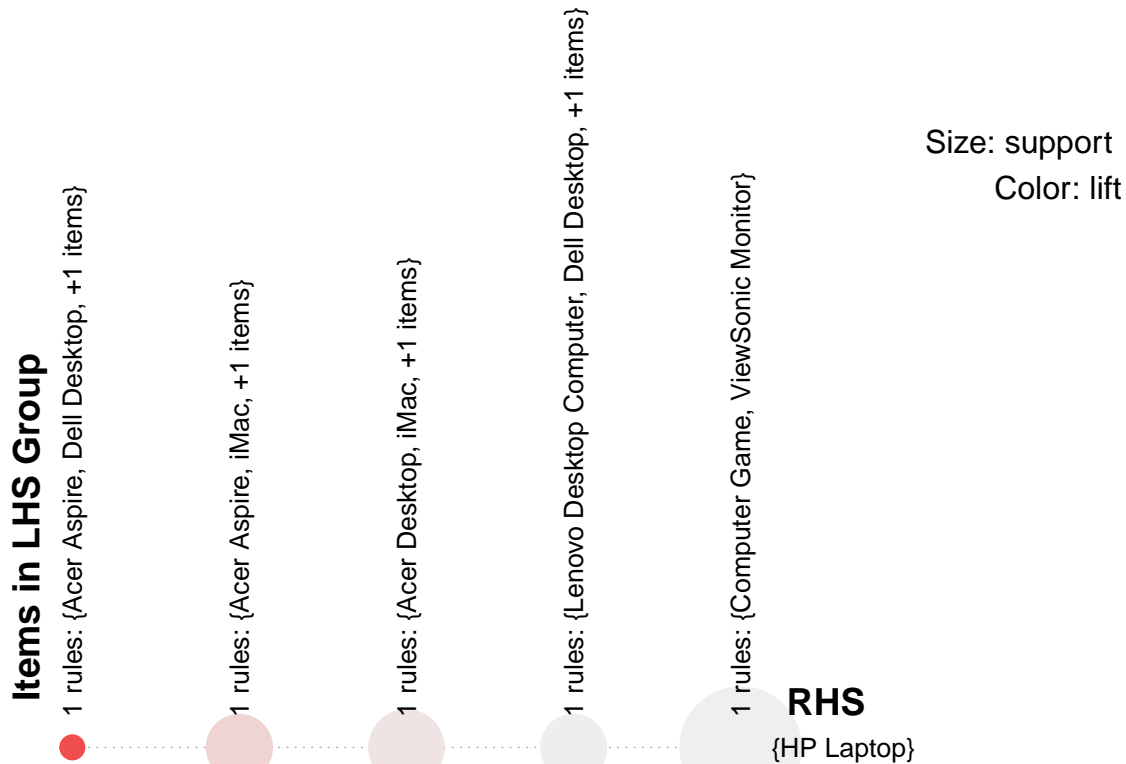
```
plot(rules1, method="paracoord", control=list(reorder=TRUE))
```

## Parallel coordinates plot for 28 rules



```
plot(topRules, method = "grouped")
```

## Grouped Matrix for 5 Rules



```
#plot(rules1,method="graph",engine='interactive',shading=NA)
```

The interactive mode performs better but can not be displayed in knit mode.

```
ItemRules <- subset(rules1, items %in% "HP Laptop")
inspect(ItemRules[1:5])
```

|        | lhs  | rhs            | support     | confidence | coverage    | lift     | count |
|--------|--|----------------|-------------|------------|-------------|----------|-------|
| ## [1] | {Acer Aspire,<br>Dell Desktop,<br>ViewSonic Monitor}             | => {HP Laptop} | 0.005287239 | 0.8125000  | 0.006507372 | 4.185928 | 52    |
| ## [2] | {Acer Aspire,<br>iMac,<br>ViewSonic Monitor}                     | => {HP Laptop} | 0.006202339 | 0.6630435  | 0.009354347 | 3.415942 | 61    |
| ## [3] | {Acer Desktop,<br>iMac,<br>ViewSonic Monitor}                    | => {HP Laptop} | 0.006405694 | 0.6363636  | 0.010066090 | 3.278489 | 63    |
| ## [4] | {Dell Desktop,<br>Lenovo Desktop Computer,<br>ViewSonic Monitor} | => {HP Laptop} | 0.006202339 | 0.6224490  | 0.009964413 | 3.206802 | 61    |
| ## [5] | {Computer Game,<br>ViewSonic Monitor}                            | => {HP Laptop} | 0.007422471 | 0.6186441  | 0.011997966 | 3.187200 | 73    |

```
rules_df <- as(rules1, "data.frame")
str(rules_df)
```

```
## 'data.frame': 28 obs. of 6 variables:
## $ rules : chr "{Acer Aspire,Dell Desktop,ViewSonic Monitor} => {HP Laptop}" "{Acer Aspire,iMac
```

```
## $ support : num 0.00529 0.0062 0.00641 0.0062 0.00742 ...
## $ confidence: num 0.812 0.663 0.636 0.622 0.619 ...
## $ coverage : num 0.00651 0.00935 0.01007 0.00996 0.012 ...
## $ lift : num 4.19 3.42 3.28 3.21 3.19 ...
## $ count : int 52 61 63 61 73 56 106 51 50 51 ...
```

```
write(rules1, file = "rules_df.csv", sep = ",", quote = TRUE, row.names = FALSE)
```

Now we know Laptop & Imax... lets understand if we can use them to purchase lower frequency products

```
rules_highest_lhs <- apriori(data= TransactionDataSet, parameter = list(supp = 0.01, conf=0.2),appearan
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
## 0.2 0.1 1 none FALSE TRUE 5 0.01 1
## maxlen target ext
## 10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
## 0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 98
##
## set item appearances ...[1 item(s)] done [0.00s].
## set transactions ...[125 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [82 item(s)] done [0.00s].
## creating transaction tree ... done [0.01s].
## checking subsets of size 1 2 done [0.00s].
## writing ... [4 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules_highest_lhs
```


```
## set of 4 rules
```

```
rules_highest_lhs <- sort(rules_highest_lhs, by ="lift")
inspect(rules_highest_lhs)
```

```
## lhs rhs support confidence coverage
## [1] {iMac} => {Dell Desktop} 0.05460092 0.2131798 0.2561261
## [2] {iMac} => {Lenovo Desktop Computer} 0.05876970 0.2294561 0.2561261
## [3] {iMac} => {HP Laptop} 0.07554652 0.2949583 0.2561261
## [4] {iMac} => {CYBERPOWER Gamer Desktop} 0.05673615 0.2215165 0.2561261
## lift count
## [1] 1.590762 537
## [2] 1.549932 578
## [3] 1.519599 743
## [4] 1.204320 558
```

```
plot(rules_highest_lhs, method="paracoord", control=list(reorder=TRUE))
```

## Parallel coordinates plot for 4 rules



```
rules_highest_lhs <- apriori(data= TransactionDataSet, parameter = list(supp = 0.03, conf=0.2),appearan
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.2    0.1    1 none FALSE          TRUE     5    0.03    1
## maxlen target  ext
##      10  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 295
##
## set item appearances ...[1 item(s)] done [0.00s].
## set transactions ...[125 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [43 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 done [0.00s].
## writing ... [6 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
rules_highest_lhs
```

```
## set of 6 rules
```

```
rules_highest_lhs <- sort(rules_highest_lhs, by = "lift")
inspect(rules_highest_lhs)
```

| ##     | lhs         | rhs                           | support    | confidence | coverage  |
|--------|-------------|-------------------------------|------------|------------|-----------|
| ## [1] | {HP Laptop} | => {ViewSonic Monitor}        | 0.04799187 | 0.2472499  | 0.1941027 |
| ## [2] | {HP Laptop} | => {Dell Desktop}             | 0.04494154 | 0.2315348  | 0.1941027 |
| ## [3] | {HP Laptop} | => {Lenovo Desktop Computer}  | 0.04616167 | 0.2378208  | 0.1941027 |
| ## [4] | {HP Laptop} | => {iMac}                     | 0.07554652 | 0.3892090  | 0.1941027 |
| ## [5] | {HP Laptop} | => {CYBERPOWER Gamer Desktop} | 0.04260295 | 0.2194866  | 0.1941027 |
| ## [6] | {}          | => {iMac}                     | 0.25612608 | 0.2561261  | 1.0000000 |

| ##     | lift     | count |
|--------|----------|-------|
| ## [1] | 2.241200 | 472   |
| ## [2] | 1.727728 | 442   |
| ## [3] | 1.606434 | 454   |
| ## [4] | 1.519599 | 743   |
| ## [5] | 1.193284 | 419   |
| ## [6] | 1.000000 | 2519  |

```
plot(rules_highest_lhs, method="paracoord", control=list(reorder=TRUE))
```

### Parallel coordinates plot for 5 rules

